

# Logistic Regression

```
In [58]: import pandas as pd
fish = pd.read_csv('dataset_Fish.csv')
fish.head()
```

```
Out[58]:
```

	Species	Weight	Length1	Length2	Length3	Height	Width
0	Bream	242.0	23.2	25.4	30.0	11.5200	4.0200
1	Bream	290.0	24.0	26.3	31.2	12.4800	4.3056
2	Bream	340.0	23.9	26.5	31.1	12.3778	4.6961
3	Bream	363.0	26.3	29.0	33.5	12.7300	4.4555
4	Bream	430.0	26.5	29.0	34.0	12.4440	5.1340

```
In [59]: fish['Species'].unique()
```

```
Out[59]: array(['Bream', 'Roach', 'Whitefish', 'Parkki', 'Perch', 'Pike', 'Smelt'],
      dtype=object)
```

```
In [60]: fish.isnull().sum()
```

```
Out[60]: Species      0
Weight      0
Length1     0
Length2     0
Length3     0
Height      0
Width       0
dtype: int64
```

```
In [61]: X = fish.iloc[:, 1:]
y = fish.loc[:, 'Species']
```

```
In [62]: X
```

```
Out[62]:
```

	Weight	Length1	Length2	Length3	Height	Width
0	242.0	23.2	25.4	30.0	11.5200	4.0200
1	290.0	24.0	26.3	31.2	12.4800	4.3056
2	340.0	23.9	26.5	31.1	12.3778	4.6961
3	363.0	26.3	29.0	33.5	12.7300	4.4555
4	430.0	26.5	29.0	34.0	12.4440	5.1340
...	...	...	...	...	...	...
154	12.2	11.5	12.2	13.4	2.0904	1.3936
155	13.4	11.7	12.4	13.5	2.4300	1.2690
156	12.2	12.1	13.0	13.8	2.2770	1.2558
157	19.7	13.2	14.3	15.2	2.8728	2.0672
158	19.9	13.8	15.0	16.2	2.9322	1.8792

159 rows × 6 columns

## Scaling the input features using MinMaxScaler

```
In [63]: y
```

```
Out[63]: 0      Bream
1      Bream
2      Bream
3      Bream
4      Bream
...
154    Smelt
155    Smelt
156    Smelt
157    Smelt
158    Smelt
Name: Species, Length: 159, dtype: object
```

```
In [65]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
scaler.fit(X)
X_scaled = scaler.transform(X)
```

## Label Encoding the target variable using LabelEncoder

```
In [76]: from sklearn.preprocessing import LabelEncoder
          label_encoder = LabelEncoder()
          y = label_encoder.fit_transform(y)
          y
```

[illegible]

## Splitting into train and test datasets using train\_test\_split

```
In [77]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

## Model Building and training

```
In [78]: from sklearn.linear_model import LogisticRegression
logReg = LogisticRegression()
# training the model
logReg.fit(X_train, y_train)
```

```
Out[78]: LogisticRegression()
```

## Predicting the output

```
In [79]: y_pred = logReg.predict(X_test)
```

## Computing the accuracy


```
In [80]: from sklearn.metrics import accuracy_score
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy: {:.2f}%".format(accuracy * 100))
```

Accuracy: 81.25%

### Confusion Matrix

```
In [50]: from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
cf = confusion_matrix(y_test, y_pred)
plt.figure()
sns.heatmap(cf, annot=True)
plt.xlabel('Prediction')
plt.ylabel('Target')
plt.title('Confusion Matrix')
```

```
Out[50]: Text(0.5, 1.0, 'Confusion Matrix')
```



## Hyperparameter tuning

```
In [82]: from sklearn.datasets import make_blobs
from sklearn.model_selection import RepeatedStratifiedKFold
from sklearn.model_selection import GridSearchCV
import warnings
warnings.filterwarnings("ignore")

# define models and parameters
solvers = ['newton-cg', 'lbfgs', 'liblinear', 'saga', 'sag']
penalty = ['l1', 'l2']
c_values = [100, 10, 1.0, 0.1, 0.01]

In [83]: # define grid search
grid = dict(solver=solvers,penalty=penalty,C=c_values)
cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
grid_search = GridSearchCV(estimator=logReg, param_grid=grid, n_jobs=-1, cv=cv, scoring='accuracy',error_score=0)
grid_result = grid_search.fit(X_train, y_train)

In [84]: # summarize results
print("Best: %f using %s\n" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
print('Mean (Std Dev) with Parameters')
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

Best: 0.942949 using {'C': 100, 'penalty': 'l1', 'solver': 'liblinear'}

Mean (Std Dev) with Parameters

```
0.000000 (0.000000) with: {'C': 100, 'penalty': 'l1', 'solver': 'newton-cg'}
0.000000 (0.000000) with: {'C': 100, 'penalty': 'l1', 'solver': 'lbfgs'}
0.942949 (0.063379) with: {'C': 100, 'penalty': 'l1', 'solver': 'liblinear'}
0.821795 (0.081394) with: {'C': 100, 'penalty': 'l1', 'solver': 'saga'}
0.000000 (0.000000) with: {'C': 100, 'penalty': 'l1', 'solver': 'sag'}
0.853419 (0.076829) with: {'C': 100, 'penalty': 'l2', 'solver': 'newton-cg'}
0.853419 (0.076829) with: {'C': 100, 'penalty': 'l2', 'solver': 'lbfgs'}
0.816667 (0.077165) with: {'C': 100, 'penalty': 'l2', 'solver': 'liblinear'}
0.808974 (0.069641) with: {'C': 100, 'penalty': 'l2', 'solver': 'saga'}
0.829701 (0.071663) with: {'C': 100, 'penalty': 'l2', 'solver': 'sag'}
0.000000 (0.000000) with: {'C': 10, 'penalty': 'l1', 'solver': 'newton-cg'}
0.000000 (0.000000) with: {'C': 10, 'penalty': 'l1', 'solver': 'lbfgs'}
0.845513 (0.088132) with: {'C': 10, 'penalty': 'l1', 'solver': 'liblinear'}
0.800855 (0.079823) with: {'C': 10, 'penalty': 'l1', 'solver': 'saga'}
0.000000 (0.000000) with: {'C': 10, 'penalty': 'l1', 'solver': 'sag'}
0.789957 (0.080672) with: {'C': 10, 'penalty': 'l2', 'solver': 'newton-cg'}
0.789957 (0.080672) with: {'C': 10, 'penalty': 'l2', 'solver': 'lbfgs'}
0.753632 (0.054284) with: {'C': 10, 'penalty': 'l2', 'solver': 'liblinear'}
0.789957 (0.080672) with: {'C': 10, 'penalty': 'l2', 'solver': 'saga'}
0.789957 (0.080672) with: {'C': 10, 'penalty': 'l2', 'solver': 'sag'}
0.000000 (0.000000) with: {'C': 1.0, 'penalty': 'l1', 'solver': 'newton-cg'}
0.000000 (0.000000) with: {'C': 1.0, 'penalty': 'l1', 'solver': 'lbfgs'}
0.700427 (0.052390) with: {'C': 1.0, 'penalty': 'l1', 'solver': 'liblinear'}
0.687393 (0.070797) with: {'C': 1.0, 'penalty': 'l1', 'solver': 'saga'}
0.000000 (0.000000) with: {'C': 1.0, 'penalty': 'l1', 'solver': 'sag'}
0.631838 (0.067975) with: {'C': 1.0, 'penalty': 'l2', 'solver': 'newton-cg'}
0.631838 (0.067975) with: {'C': 1.0, 'penalty': 'l2', 'solver': 'lbfgs'}
0.610684 (0.061014) with: {'C': 1.0, 'penalty': 'l2', 'solver': 'liblinear'}
0.631838 (0.067975) with: {'C': 1.0, 'penalty': 'l2', 'solver': 'saga'}
0.631838 (0.067975) with: {'C': 1.0, 'penalty': 'l2', 'solver': 'sag'}
0.000000 (0.000000) with: {'C': 0.1, 'penalty': 'l1', 'solver': 'newton-cg'}
0.000000 (0.000000) with: {'C': 0.1, 'penalty': 'l1', 'solver': 'lbfgs'}
0.369231 (0.023500) with: {'C': 0.1, 'penalty': 'l1', 'solver': 'liblinear'}
0.369231 (0.023500) with: {'C': 0.1, 'penalty': 'l1', 'solver': 'saga'}
0.000000 (0.000000) with: {'C': 0.1, 'penalty': 'l1', 'solver': 'sag'}
0.369231 (0.023500) with: {'C': 0.1, 'penalty': 'l2', 'solver': 'newton-cg'}
0.369231 (0.023500) with: {'C': 0.1, 'penalty': 'l2', 'solver': 'lbfgs'}
0.369231 (0.023500) with: {'C': 0.1, 'penalty': 'l2', 'solver': 'liblinear'}
0.369231 (0.023500) with: {'C': 0.1, 'penalty': 'l2', 'solver': 'saga'}
0.369231 (0.023500) with: {'C': 0.1, 'penalty': 'l2', 'solver': 'sag'}
0.000000 (0.000000) with: {'C': 0.01, 'penalty': 'l1', 'solver': 'newton-cg'}
0.000000 (0.000000) with: {'C': 0.01, 'penalty': 'l1', 'solver': 'lbfgs'}
0.196154 (0.034899) with: {'C': 0.01, 'penalty': 'l1', 'solver': 'liblinear'}
0.358547 (0.048929) with: {'C': 0.01, 'penalty': 'l1', 'solver': 'saga'}
0.000000 (0.000000) with: {'C': 0.01, 'penalty': 'l1', 'solver': 'sag'}
0.369231 (0.023500) with: {'C': 0.01, 'penalty': 'l2', 'solver': 'newton-cg'}
0.369231 (0.023500) with: {'C': 0.01, 'penalty': 'l2', 'solver': 'lbfgs'}
0.369231 (0.023500) with: {'C': 0.01, 'penalty': 'l2', 'solver': 'liblinear'}
0.369231 (0.023500) with: {'C': 0.01, 'penalty': 'l2', 'solver': 'saga'}
0.369231 (0.023500) with: {'C': 0.01, 'penalty': 'l2', 'solver': 'sag'}
```

## SVC\_Apple\_orange


```
In [1]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import confusion_matrix
from sklearn.preprocessing import LabelEncoder
```

```
In [4]: data = pd.read_csv('apples_and_oranges.csv')
print(data)
```

	Weight	Size	Class
0	69	4.39	orange
1	69	4.21	orange
2	65	4.09	orange
3	72	5.85	apple
4	67	4.70	orange
5	73	5.68	apple
6	70	5.56	apple
7	75	5.11	apple
8	74	5.36	apple
9	65	4.27	orange
10	73	5.79	apple
11	70	5.47	apple
12	74	5.53	apple
13	68	4.47	orange
14	74	5.22	apple
15	65	4.48	orange
16	69	4.66	orange
17	75	5.25	apple
18	67	4.18	orange
19	74	5.50	apple
20	66	4.13	orange
21	70	4.83	orange
22	69	4.61	orange
23	68	4.08	orange
24	67	4.25	orange
25	71	5.35	apple
26	67	4.01	orange
27	70	4.22	orange
28	74	5.25	apple
29	71	5.26	apple
30	73	5.78	apple
31	66	4.68	orange
32	72	5.72	apple
33	73	5.17	apple
34	68	4.83	orange
35	69	4.11	orange
36	69	4.76	orange
37	74	5.48	apple
38	70	5.59	apple
39	73	5.03	apple

```
In [5]: import seaborn as sns
sns.scatterplot(x="Weight",
               y="Size",
               hue="Class",
               data=data)
```

```
Out[5]: <AxesSubplot:xlabel='Weight', ylabel='Size'>
```



```
In [6]: # splitting data into training and test set
training_set, test_set = train_test_split(data, test_size=0.2, random_state=1)
print("train:", training_set)
print("test:", test_set)
```

train:	Weight	Size	Class
19	74	5.50	apple
26	67	4.01	orange
32	72	5.72	apple
17	75	5.25	apple
30	73	5.78	apple
36	69	4.76	orange
33	73	5.17	apple
28	74	5.25	apple
4	67	4.70	orange
14	74	5.22	apple
10	73	5.79	apple
35	69	4.11	orange
23	68	4.08	orange
24	67	4.25	orange
34	68	4.83	orange
20	66	4.13	orange
18	67	4.18	orange
25	71	5.35	apple
6	70	5.56	apple
13	68	4.47	orange
7	75	5.11	apple
38	70	5.59	apple
1	69	4.21	orange
16	69	4.66	orange
0	69	4.39	orange
15	65	4.48	orange
5	73	5.68	apple
11	70	5.47	apple
9	65	4.27	orange
8	74	5.36	apple
12	74	5.53	apple
37	74	5.48	apple
test:	Weight	Size	Class
2	65	4.09	orange
31	66	4.68	orange
3	72	5.85	apple
21	70	4.83	orange
27	70	4.22	orange
29	71	5.26	apple
22	69	4.61	orange
39	73	5.03	apple

```
In [7]: x_train = training_set.iloc[:,0:2].values # data
y_train = training_set.iloc[:,2].values # target
x_test = test_set.iloc[:,0:2].values # data
y_test = test_set.iloc[:,2].values # target
print(x_train,y_train)
print(x_test,y_test)

[[74.    5.5 ]
 [67.    4.01]
 [72.    5.72]
 [75.    5.25]
 [73.    5.78]
 [69.    4.76]
 [73.    5.17]
 [74.    5.25]
 [67.    4.7 ]
 [74.    5.22]
 [73.    5.79]
 [69.    4.11]
 [68.    4.08]
 [67.    4.25]
 [68.    4.83]
 [66.    4.13]
 [67.    4.18]
 [71.    5.35]
 [70.    5.56]
 [68.    4.47]
 [75.    5.11]
 [70.    5.59]
 [69.    4.21]
 [69.    4.66]
 [69.    4.39]
 [65.    4.48]
 [73.    5.68]
 [70.    5.47]
 [65.    4.27]
 [74.    5.36]
 [74.    5.53]
 [74.    5.48]] ['apple' 'orange' 'apple' 'apple' 'apple' 'orange' 'apple' 'apple'
 'orange' 'apple' 'apple' 'orange' 'orange' 'orange' 'orange' 'orange'
 'orange' 'apple' 'apple' 'orange' 'apple' 'apple' 'orange' 'orange'
 'orange' 'orange' 'apple' 'apple' 'orange' 'apple' 'apple' 'apple']
[[65.    4.09]
 [66.    4.68]
 [72.    5.85]
 [70.    4.83]
 [70.    4.22]
 [71.    5.26]
 [69.    4.61]
 [73.    5.03]] ['orange' 'orange' 'apple' 'orange' 'orange' 'apple' 'orange' 'apple']
```

```
In [8]: # fitting the data (train a model)
classifier = SVC(kernel='rbf',random_state=1,C=1,gamma='auto')
classifier.fit(x_train,y_train)
```

```
Out[8]: SVC
SVC(C=1, gamma='auto', random_state=1)
```

```
In [9]: y_pred = classifier.predict(x_test)
print(y_pred)

['orange' 'orange' 'apple' 'apple' 'orange' 'apple' 'orange' 'apple']
```

```
In [ ]: # creating confusion matrix and accuracy calculation
cm = confusion_matrix(y_test,y_pred)
print(cm)
accuracy = float(cm.diagonal().sum())/len(y_test)
print('model accuracy is:',accuracy*100,'%')

#x1_test = [[73,6]] # for new data testing
```